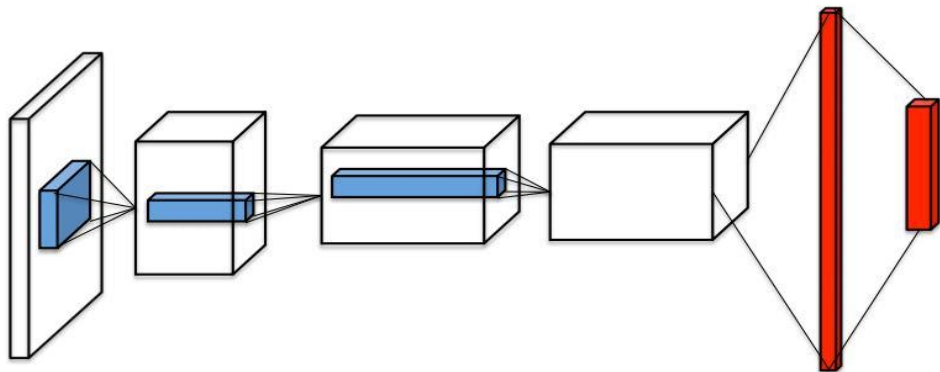


Duelling Networks

Nan Rosemary Ke

Current DQN architecture

- Uses a single network for computing the q-value updates, given a state s and action a (s, a).
- Problems:
 - Existing deep neural network architectures may not be suited for Reinforcement Learning
 - The state value function and the advantage function updates at different rates.



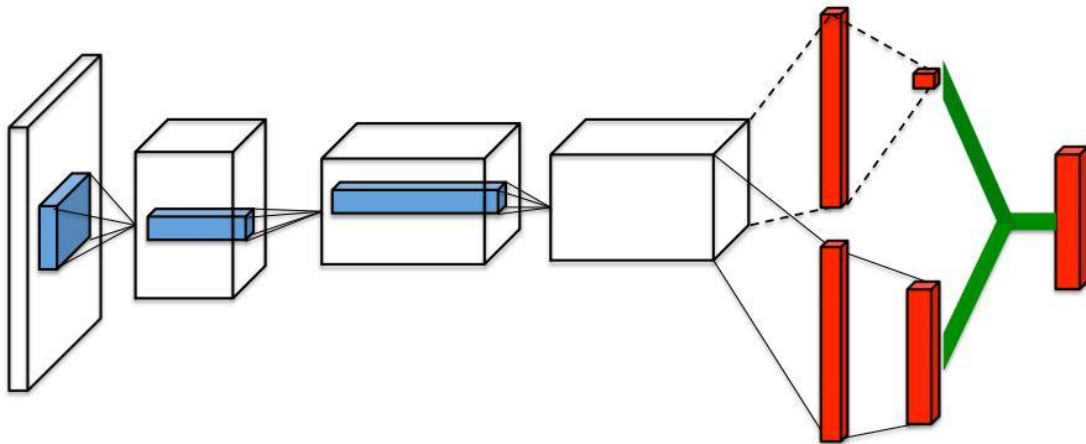
Motivation

- Novel Deep Neural Network architectures build for Model-free RL problems.
- Having separate network approximating the State-value function approximator and the advantage update function.
 - State-value function and advantage-value function updates at different pace.



Duelling Network architecture

- Shared Convolutional feature learning module. Separate state value and advantage function approximator.
 - No extra supervision.
 - Learns which state is valuable, do not have to learn action-value for all action for all states



Duelling network

Definitions

- State value $V(s)$
- Advantage function $A(s, a)$
- Policy π
- Return $R_t = \sum_{\tau=t}^{\infty} \gamma^{\tau-t} r_{\tau}, \gamma \in [0, 1]$
- Q function $Q^{\pi}(s, a) = E[R_t | S_t = s, a_t = a, \pi]$
- State-value $V^{\pi}(s) = E_{a \sim \pi(s)}[Q^{\pi}(s, a)]$



Bellman equation vs Advantage equation

- Bellman equation

$$Q^*(s, a) = \mathbb{E}_{s'} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

- Advantage equation

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

- $\mathbb{E}_{a \sim \pi(s')} [A^\pi(s, a)] = 0$



Duelling network - intuition

- Value $V(s)$ computes **how good it is to be in that particular state**
- $Q(s,a)$ computes, given in state s , how **good is action a** .
- $A = V - Q$ computes the **relative importance of each action**

Duelling network separates the computation of the value of the advantage function, since they update at different rates.



Duelling network - Key insights

- For many states
 - Not necessary to compute all action values.
 - In many states, the action has no effect.
- Bootstrapping algorithm
 - Computation of state value is extremely important
 - Bootstrapping: updating estimates based on other estimates (other estimate should be fairly accurate)



Duelling network - Formulation

- $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$
- $V^\pi(s) = \mathbb{E}_{a \sim \pi(s)}[Q^\pi(s, a)]$
 - $A^\pi(s, a) = Q^\pi(s, a) - \mathbb{E}_{a \sim \pi(s)}[Q^\pi(s, a)]$
- $\mathbb{E}_{a \sim \pi(s)}[A^\pi(s, a)] = 0$
- For a deterministic policy, $a^* = \arg\max_{a' \in \mathcal{A}} Q(s, a')$
 - $Q(s, a^*) = V(s)$ and $A(s, a^*) = 0$



Duelling network - formulation


- Architecture
 - CNN + fully connected layers that output
 - Scalar $V(s)$
 - Vector $A(s, a)$
- Tempt to construct aggregation module
 - $Q(s, a; \theta, \alpha, \beta) = V(s; \theta, \beta) + A(s, a; \theta, \alpha)$



Results in the paper


When new car is

Coming, focus on horizon/



Don't pay attention
when there is no car

Focus on the score



Attention on the car immediately
at the front. Making its choice of
actions relevant.

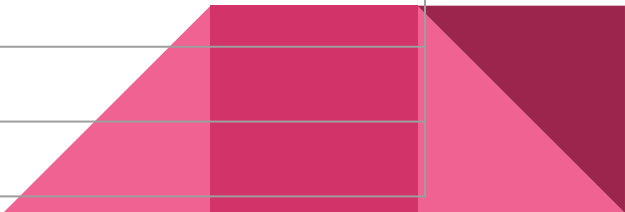
My Experiments

- Architecture
 - 2 feedforward dense layers (layer size 100)
 - Separate Q and A layer, merges into a single action-value function estimator
- Optimizer:
 - Adam
 - Learning rate: 0.001
- Game : Carpool



My Experiments

Espisodes	Time steps	Reward
1	18	18
2	28	28
3	48	48
4	200	200
5	200	200
6	200	200
7	200	200
8	200	200
9	200	200



My Experiments

Carpole - Reward over episodes

